

Machine Learning Laboratory

(410302)

BE Sem I Honors in AI/ML

Academic Year: 2021-22

Mini Project

Name: Aboli Marathe

Roll Number: 41301

Branch: Department of Computer Engineering

▼ Problem Statement

Introduction

Reinforcement learning is an area of machine learning concerned with how intelligent agents ought to take actions in an environment in order to maximize the notion of cumulative reward.

Cartpole problem

Cartpole - known also as an Inverted Pendulum is a pendulum with a center of gravity above its pivot point. It's unstable, but can be controlled by moving the pivot point under the center of mass. The goal is to keep the cartpole balanced by applying appropriate forces to a pivot point.

In this Mini project, we will implement a deep reinforcement learning algorithms for solving the cartpole problem and analyze the performance of the model.

Problem Statement: Evaluating the performance of Deep Reinforcement Learning for solving the Cartpole Problem

Algorithm Employed: Rainbow

A selection of the below six extensions that each have addressed a limitation and improved overall performance were integrated into a single integrated agent, Rainbow for improved performance on RL tasks [8]:

1. Double Q-learning
2. Prioritized replay
3. Dueling networks
4. Multi-step learning
5. Distributional RL
6. Noisy Nets

1. V. Mnih et al., "Human-level control through deep reinforcement learning." Nature, 518 (7540):529–533, 2015.
2. van Hasselt et al., "Deep Reinforcement Learning with Double Q-learning." arXiv preprint arXiv:1509.06461, 2015.
3. T. Schaul et al., "Prioritized Experience Replay." arXiv preprint arXiv:1511.05952, 2015.
4. Z. Wang et al., "Dueling Network Architectures for Deep Reinforcement Learning." arXiv preprint arXiv:1511.06581, 2015.
5. M. Fortunato et al., "Noisy Networks for Exploration." arXiv preprint arXiv:1706.10295, 2017.
6. M. G. Bellemare et al., "A Distributional Perspective on Reinforcement Learning." arXiv preprint arXiv:1707.06887, 2017.
7. R. S. Sutton, "Learning to predict by the methods of temporal differences." Machine learning, 3(1):9–44, 1988.
8. M. Hessel et al., "Rainbow: Combining Improvements in Deep Reinforcement Learning." arXiv preprint arXiv:1710.02298, 2017.

▼ Configurations for Colab

```
import sys
IN_COLAB = "google.colab" in sys.modules

if IN_COLAB:
    !apt install python-opengl
    !apt install ffmpeg
    !apt install xvfb
    !pip install pyvirtualdisplay
    !pip install gym
    from pyvirtualdisplay import Display

    # Start virtual display
    dis = Display(visible=0, size=(400, 400))
    dis.start()

    Reading package lists... Done
    Building dependency tree
    Reading state information... Done
    python-opengl is already the newest version (3.1.0+dfsg-1).
    The following package was automatically installed and is no longer required:
```

```

libnvidia-common-460
Use 'apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 37 not upgraded.
Reading package lists... Done
Building dependency tree
Reading state information... Done
ffmpeg is already the newest version (7:3.4.8-0ubuntu0.2).
The following package was automatically installed and is no longer required:
  libnvidia-common-460
Use 'apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 37 not upgraded.
Reading package lists... Done
Building dependency tree
Reading state information... Done
xvfb is already the newest version (2:1.19.6-1ubuntu4.9).
The following package was automatically installed and is no longer required:
  libnvidia-common-460
Use 'apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 37 not upgraded.
Requirement already satisfied: pyvirtualdisplay in /usr/local/lib/python3.7/dist-packa
Requirement already satisfied: EasyProcess in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: gym in /usr/local/lib/python3.7/dist-packages (0.17.3)
Requirement already satisfied: cloudpickle<1.7.0,>=1.2.0 in /usr/local/lib/python3.7/
Requirement already satisfied: numpy>=1.10.4 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: pygame<=1.5.0,>=1.4.0 in /usr/local/lib/python3.7/dist
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from

```


▼ 08. Rainbow

[M. Hessel et al., "Rainbow: Combining Improvements in Deep Reinforcement Learning." arXiv preprint arXiv:1710.02298, 2017.](https://arxiv.org/abs/1710.02298)

We will integrate all the following seven components into a single integrated agent, which is called Rainbow!

1. DQN
2. Double DQN
3. Prioritized Experience Replay
4. Dueling Network
5. Noisy Network
6. Categorical DQN
7. N-step Learning

This method shows an impressive performance on the Atari 2600 benchmark, both in terms of data efficiency and final performance.



However, the integration is not so simple because some of components are not independent each other, so we will look into a number of points that people especially feel confused.

1. Noisy Network <-> Dueling Network
2. Dueling Network <-> Categorical DQN
3. Categorical DQN <-> Double DQN

```
import math
import os
import random
from collections import deque
from typing import Deque, Dict, List, Tuple

import gym
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from IPython.display import clear_output
from torch.nn.utils import clip_grad_norm_

# download segment tree module
if IN_COLAB:
    !wget https://raw.githubusercontent.com/curt-park/rainbow-is-all-you-need/master/segme

from segment_tree import MinSegmentTree, SumSegmentTree
```

--2021-10-30 07:44:42-- https://raw.githubusercontent.com/curt-park/rainbow-is-all-you-need/master/segment_tree.py

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 1
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|
HTTP request sent, awaiting response... 200 OK
Length: 4283 (4.2K) [text/plain]
Saving to: 'segment_tree.py.1'
```

```
segment_tree.py.1 100%[=====>] 4.18K --.-KB/s in 0s
```

```
2021-10-30 07:44:42 (46.4 MB/s) - 'segment_tree.py.1' saved [4283/4283]
```



▼ Replay buffer

Same as the basic N-step buffer.

(Please see *01.dqn.ipynb*, *07.n_step_learning.ipynb* for detailed description about the basic (n-step) replay buffer.)

```
class ReplayBuffer:
    """A simple numpy replay buffer."""

    def __init__(
        self,
        obs_dim: int,
        size: int,
        batch_size: int = 32,
        n_step: int = 1,
        gamma: float = 0.99
    ):
        self.obs_buf = np.zeros([size, obs_dim], dtype=np.float32)
        self.next_obs_buf = np.zeros([size, obs_dim], dtype=np.float32)
        self.acts_buf = np.zeros([size], dtype=np.float32)
        self.rews_buf = np.zeros([size], dtype=np.float32)
        self.done_buf = np.zeros(size, dtype=np.float32)
        self.max_size, self.batch_size = size, batch_size
        self.ptr, self.size, = 0, 0

        # for N-step Learning
        self.n_step_buffer = deque(maxlen=n_step)
        self.n_step = n_step
        self.gamma = gamma

    def store(
        self,
        obs: np.ndarray,
        act: np.ndarray,
        rew: float,
        next_obs: np.ndarray,
        done: bool,
    ) -> Tuple[np.ndarray, np.ndarray, float, np.ndarray, bool]:
        transition = (obs, act, rew, next_obs, done)
        self.n_step_buffer.append(transition)
```

```

# single step transition is not ready
if len(self.n_step_buffer) < self.n_step:
    return ()

# make a n-step transition
rew, next_obs, done = self._get_n_step_info(
    self.n_step_buffer, self.gamma
)
obs, act = self.n_step_buffer[0][:2]

self.obs_buf[self.ptr] = obs
self.next_obs_buf[self.ptr] = next_obs
self.acts_buf[self.ptr] = act
self.rews_buf[self.ptr] = rew
self.done_buf[self.ptr] = done
self.ptr = (self.ptr + 1) % self.max_size
self.size = min(self.size + 1, self.max_size)

return self.n_step_buffer[0]

def sample_batch(self) -> Dict[str, np.ndarray]:
    idxs = np.random.choice(self.size, size=self.batch_size, replace=False)

    return dict(
        obs=self.obs_buf[idxs],
        next_obs=self.next_obs_buf[idxs],
        acts=self.acts_buf[idxs],
        rews=self.rews_buf[idxs],
        done=self.done_buf[idxs],
        # for N-step Learning
        indices=idxs,
    )

def sample_batch_from_idx(
    self, idxs: np.ndarray
) -> Dict[str, np.ndarray]:
    # for N-step Learning
    return dict(
        obs=self.obs_buf[idxs],
        next_obs=self.next_obs_buf[idxs],
        acts=self.acts_buf[idxs],
        rews=self.rews_buf[idxs],
        done=self.done_buf[idxs],
    )

def _get_n_step_info(
    self, n_step_buffer: Deque, gamma: float
) -> Tuple[np.int64, np.ndarray, bool]:
    """Return n step rew, next_obs, and done."""
    # info of the last transition
    rew, next_obs, done = n_step_buffer[-1][-3:]

    for transition in reversed(list(n_step_buffer[:-1])):
        r, n_o, d = transition[-3:]

```

```

        rew = r + gamma * rew * (1 - d)
        next_obs, done = (n_o, d) if d else (next_obs, done)

    return rew, next_obs, done

def __len__(self) -> int:
    return self.size

```

▼ Prioritized replay Buffer

store method returns boolean in order to inform if a N-step transition has been generated.

(Please see *02.per.ipynb* for detailed description about PER.)

```

class PrioritizedReplayBuffer(ReplayBuffer):
    """Prioritized Replay buffer.

    Attributes:
        max_priority (float): max priority
        tree_ptr (int): next index of tree
        alpha (float): alpha parameter for prioritized replay buffer
        sum_tree (SumSegmentTree): sum tree for prior
        min_tree (MinSegmentTree): min tree for min prior to get max weight

    """

    def __init__(
        self,
        obs_dim: int,
        size: int,
        batch_size: int = 32,
        alpha: float = 0.6,
        n_step: int = 1,
        gamma: float = 0.99,
    ):
        """Initialization."""
        assert alpha >= 0

        super(PrioritizedReplayBuffer, self).__init__(
            obs_dim, size, batch_size, n_step, gamma
        )
        self.max_priority, self.tree_ptr = 1.0, 0
        self.alpha = alpha

        # capacity must be positive and a power of 2.
        tree_capacity = 1
        while tree_capacity < self.max_size:
            tree_capacity *= 2

        self.sum_tree = SumSegmentTree(tree_capacity)
        self.min_tree = MinSegmentTree(tree_capacity)

```

```

def store(
    self,
    obs: np.ndarray,
    act: int,
    rew: float,
    next_obs: np.ndarray,
    done: bool,
) -> Tuple[np.ndarray, np.ndarray, float, np.ndarray, bool]:
    """Store experience and priority."""
    transition = super().store(obs, act, rew, next_obs, done)

    if transition:
        self.sum_tree[self.tree_ptr] = self.max_priority ** self.alpha
        self.min_tree[self.tree_ptr] = self.max_priority ** self.alpha
        self.tree_ptr = (self.tree_ptr + 1) % self.max_size

    return transition

def sample_batch(self, beta: float = 0.4) -> Dict[str, np.ndarray]:
    """Sample a batch of experiences."""
    assert len(self) >= self.batch_size
    assert beta > 0

    indices = self._sample_proportional()

    obs = self.obs_buf[indices]
    next_obs = self.next_obs_buf[indices]
    acts = self.acts_buf[indices]
    rews = self.rews_buf[indices]
    done = self.done_buf[indices]
    weights = np.array([self._calculate_weight(i, beta) for i in indices])

    return dict(
        obs=obs,
        next_obs=next_obs,
        acts=acts,
        rews=rews,
        done=done,
        weights=weights,
        indices=indices,
    )

def update_priorities(self, indices: List[int], priorities: np.ndarray):
    """Update priorities of sampled transitions."""
    assert len(indices) == len(priorities)

    for idx, priority in zip(indices, priorities):
        assert priority > 0
        assert 0 <= idx < len(self)

        self.sum_tree[idx] = priority ** self.alpha
        self.min_tree[idx] = priority ** self.alpha

        self.max_priority = max(self.max_priority, priority)

```



```

def _sample_proportional(self) -> List[int]:
    """Sample indices based on proportions."""
    indices = []
    p_total = self.sum_tree.sum(0, len(self) - 1)
    segment = p_total / self.batch_size

    for i in range(self.batch_size):
        a = segment * i
        b = segment * (i + 1)
        upperbound = random.uniform(a, b)
        idx = self.sum_tree.retrieve(upperbound)
        indices.append(idx)

    return indices

def _calculate_weight(self, idx: int, beta: float):
    """Calculate the weight of the experience at idx."""
    # get max weight
    p_min = self.min_tree.min() / self.sum_tree.sum()
    max_weight = (p_min * len(self)) ** (-beta)

    # calculate weights
    p_sample = self.sum_tree[idx] / self.sum_tree.sum()
    weight = (p_sample * len(self)) ** (-beta)
    weight = weight / max_weight

    return weight

```

▼ Noisy Layer

Please see *05.noisy_net.ipynb* for detailed description.

References:

- <https://github.com/higgsfield/RL-Adventure/blob/master/5.noisy%20dqn.ipynb>
- <https://github.com/Kaixhin/Rainbow/blob/master/model.py>

```

class NoisyLinear(nn.Module):
    """Noisy linear module for NoisyNet.

```

Attributes:

```

    in_features (int): input size of linear module
    out_features (int): output size of linear module
    std_init (float): initial std value
    weight_mu (nn.Parameter): mean value weight parameter
    weight_sigma (nn.Parameter): std value weight parameter
    bias_mu (nn.Parameter): mean value bias parameter
    bias_sigma (nn.Parameter): std value bias parameter

```

```

    """

```

```

def __init__(
    self,
    in_features: int,
    out_features: int,
    std_init: float = 0.5,
):
    """Initialization."""
    super(NoisyLinear, self).__init__()

    self.in_features = in_features
    self.out_features = out_features
    self.std_init = std_init

    self.weight_mu = nn.Parameter(torch.Tensor(out_features, in_features))
    self.weight_sigma = nn.Parameter(
        torch.Tensor(out_features, in_features)
    )
    self.register_buffer(
        "weight_epsilon", torch.Tensor(out_features, in_features)
    )

    self.bias_mu = nn.Parameter(torch.Tensor(out_features))
    self.bias_sigma = nn.Parameter(torch.Tensor(out_features))
    self.register_buffer("bias_epsilon", torch.Tensor(out_features))

    self.reset_parameters()
    self.reset_noise()

def reset_parameters(self):
    """Reset trainable network parameters (factorized gaussian noise)."""
    mu_range = 1 / math.sqrt(self.in_features)
    self.weight_mu.data.uniform_(-mu_range, mu_range)
    self.weight_sigma.data.fill_(
        self.std_init / math.sqrt(self.in_features)
    )
    self.bias_mu.data.uniform_(-mu_range, mu_range)
    self.bias_sigma.data.fill_(
        self.std_init / math.sqrt(self.out_features)
    )

def reset_noise(self):
    """Make new noise."""
    epsilon_in = self.scale_noise(self.in_features)
    epsilon_out = self.scale_noise(self.out_features)

    # outer product
    self.weight_epsilon.copy_(epsilon_out.ger(epsilon_in))
    self.bias_epsilon.copy_(epsilon_out)

def forward(self, x: torch.Tensor) -> torch.Tensor:
    """Forward method implementation.

    We don't use separate statements on train / eval mode.
    It doesn't show remarkable difference of performance.
    """

```

```

        return F.linear(
            x,
            self.weight_mu + self.weight_sigma * self.weight_epsilon,
            self.bias_mu + self.bias_sigma * self.bias_epsilon,
        )

    @staticmethod
    def scale_noise(size: int) -> torch.Tensor:
        """Set scale to make noise (factorized gaussian noise)."""
        x = torch.randn(size)

        return x.sign().mul(x.abs().sqrt())

```

▼ NoisyNet + DuelingNet + Categorical DQN

NoisyNet + DuelingNet

NoisyLinear is employed for the last two layers of advantage and value layers. The noise should be reset at every update step.

DuelingNet + Categorical DQN

The dueling network architecture is adapted for use with return distributions. The network has a shared representation, which is then fed into a value stream with `atom_size` outputs, and into an advantage stream with `atom_size × out_dim` outputs. For each atom, the value and advantage streams are aggregated, as in dueling DQN, and then passed through a softmax layer to obtain the normalized parametric distributions used to estimate the returns' distributions.

```

advantage = self.advantage_layer(adv_hid).view(-1, self.out_dim, self.atom_size)
value = self.value_layer(val_hid).view(-1, 1, self.atom_size)
q_atoms = value + advantage - advantage.mean(dim=1, keepdim=True)

dist = F.softmax(q_atoms, dim=-1)

```

(Please see *04.dueling.ipynb*, *05.noisy_net.ipynb*, *06.categorical_dqn.ipynb* for detailed description of each component's network architecture.)

```

class Network(nn.Module):
    def __init__(
        self,
        in_dim: int,
        out_dim: int,
        atom_size: int,
        support: torch.Tensor
    ):
        """Initialization."""
        super(Network, self).__init__()

        self.support = support

```

```

self.out_dim = out_dim
self.atom_size = atom_size

# set common feature layer
self.feature_layer = nn.Sequential(
    nn.Linear(in_dim, 128),
    nn.ReLU(),
)

# set advantage layer
self.advantage_hidden_layer = NoisyLinear(128, 128)
self.advantage_layer = NoisyLinear(128, out_dim * atom_size)

# set value layer
self.value_hidden_layer = NoisyLinear(128, 128)
self.value_layer = NoisyLinear(128, atom_size)

def forward(self, x: torch.Tensor) -> torch.Tensor:
    """Forward method implementation."""
    dist = self.dist(x)
    q = torch.sum(dist * self.support, dim=2)

    return q

def dist(self, x: torch.Tensor) -> torch.Tensor:
    """Get distribution for atoms."""
    feature = self.feature_layer(x)
    adv_hid = F.relu(self.advantage_hidden_layer(feature))
    val_hid = F.relu(self.value_hidden_layer(feature))

    advantage = self.advantage_layer(adv_hid).view(
        -1, self.out_dim, self.atom_size
    )
    value = self.value_layer(val_hid).view(-1, 1, self.atom_size)
    q_atoms = value + advantage - advantage.mean(dim=1, keepdim=True)

    dist = F.softmax(q_atoms, dim=-1)
    dist = dist.clamp(min=1e-3) # for avoiding nans

    return dist

def reset_noise(self):
    """Reset all noisy layers."""
    self.advantage_hidden_layer.reset_noise()
    self.advantage_layer.reset_noise()
    self.value_hidden_layer.reset_noise()
    self.value_layer.reset_noise()

```

▼ Rainbow Agent

Here is a summary of DQNAgent class.

Method	Note
select_action	select an action from the input state.

Method	Note
step	take an action and return the response of the env.
compute_dqn_loss	return dqn loss.
update_model	update the model by gradient descent.
target_hard_update	hard update from the local model to the target model.
train	train the agent during num_frames.
test	test the agent (1 episode).
plot	plot the training progresses.

Categorical DQN + Double DQN

The idea of Double Q-learning is to reduce overestimations by decomposing the max operation in the target into action selection and action evaluation. Here, we use `self.dqn` instead of `self.dqn_target` to obtain the target actions.

```
# Categorical DQN + Double DQN
# target_dqn is used when we don't employ double DQN
next_action = self.dqn(next_state).argmax(1)
next_dist = self.dqn_target.dist(next_state)
next_dist = next_dist[range(self.batch_size), next_action]
```

```
class DQNAgent:
    """DQN Agent interacting with environment.

    Attribute:
        env (gym.Env): openAI Gym environment
        memory (PrioritizedReplayBuffer): replay memory to store transitions
        batch_size (int): batch size for sampling
        target_update (int): period for target model's hard update
        gamma (float): discount factor
        dqn (Network): model to train and select actions
        dqn_target (Network): target model to update
        optimizer (torch.optim): optimizer for training dqn
        transition (list): transition information including
            state, action, reward, next_state, done
        v_min (float): min value of support
        v_max (float): max value of support
        atom_size (int): the unit number of support
        support (torch.Tensor): support for categorical dqn
        use_n_step (bool): whether to use n_step memory
        n_step (int): step number to calculate n-step td error
        memory_n (ReplayBuffer): n-step replay buffer
    """

    def __init__(
        self,
        env: gym.Env,
        memory_size: int,
```

```

    batch_size: int,
    target_update: int,
    gamma: float = 0.99,
    # PER parameters
    alpha: float = 0.2,
    beta: float = 0.6,
    prior_eps: float = 1e-6,
    # Categorical DQN parameters
    v_min: float = 0.0,
    v_max: float = 200.0,
    atom_size: int = 51,
    # N-step Learning
    n_step: int = 3,
):
    """Initialization.

    Args:
        env (gym.Env): openAI Gym environment
        memory_size (int): length of memory
        batch_size (int): batch size for sampling
        target_update (int): period for target model's hard update
        lr (float): learning rate
        gamma (float): discount factor
        alpha (float): determines how much prioritization is used
        beta (float): determines how much importance sampling is used
        prior_eps (float): guarantees every transition can be sampled
        v_min (float): min value of support
        v_max (float): max value of support
        atom_size (int): the unit number of support
        n_step (int): step number to calculate n-step td error
    """
    obs_dim = env.observation_space.shape[0]
    action_dim = env.action_space.n

    self.env = env
    self.batch_size = batch_size
    self.target_update = target_update
    self.gamma = gamma
    # NoisyNet: All attributes related to epsilon are removed

    # device: cpu / gpu
    self.device = torch.device(
        "cuda" if torch.cuda.is_available() else "cpu"
    )
    print(self.device)

    # PER
    # memory for 1-step Learning
    self.beta = beta
    self.prior_eps = prior_eps
    self.memory = PrioritizedReplayBuffer(
        obs_dim, memory_size, batch_size, alpha=alpha
    )

    # memory for N-step Learning

```

```

self.use_n_step = True if n_step > 1 else False
if self.use_n_step:
    self.n_step = n_step
    self.memory_n = ReplayBuffer(
        obs_dim, memory_size, batch_size, n_step=n_step, gamma=gamma
    )

# Categorical DQN parameters
self.v_min = v_min
self.v_max = v_max
self.atom_size = atom_size
self.support = torch.linspace(
    self.v_min, self.v_max, self.atom_size
).to(self.device)

# networks: dqn, dqn_target
self.dqn = Network(
    obs_dim, action_dim, self.atom_size, self.support
).to(self.device)
self.dqn_target = Network(
    obs_dim, action_dim, self.atom_size, self.support
).to(self.device)
self.dqn_target.load_state_dict(self.dqn.state_dict())
self.dqn_target.eval()

# optimizer
self.optimizer = optim.Adam(self.dqn.parameters())

# transition to store in memory
self.transition = list()

# mode: train / test
self.is_test = False

def select_action(self, state: np.ndarray) -> np.ndarray:
    """Select an action from the input state."""
    # NoisyNet: no epsilon greedy action selection
    selected_action = self.dqn(
        torch.FloatTensor(state).to(self.device)
    ).argmax()
    selected_action = selected_action.detach().cpu().numpy()

    if not self.is_test:
        self.transition = [state, selected_action]

    return selected_action

def step(self, action: np.ndarray) -> Tuple[np.ndarray, np.float64, bool]:
    """Take an action and return the response of the env."""
    next_state, reward, done, _ = self.env.step(action)

    if not self.is_test:
        self.transition += [reward, next_state, done]

    # N-step transition

```

```

        if self.use_n_step:
            one_step_transition = self.memory_n.store(*self.transition)
            # 1-step transition
        else:
            one_step_transition = self.transition

        # add a single step transition
        if one_step_transition:
            self.memory.store(*one_step_transition)

    return next_state, reward, done

def update_model(self) -> torch.Tensor:
    """Update the model by gradient descent."""
    # PER needs beta to calculate weights
    samples = self.memory.sample_batch(self.beta)
    weights = torch.FloatTensor(
        samples["weights"].reshape(-1, 1)
    ).to(self.device)
    indices = samples["indices"]

    # 1-step Learning loss
    elementwise_loss = self._compute_dqn_loss(samples, self.gamma)

    # PER: importance sampling before average
    loss = torch.mean(elementwise_loss * weights)

    # N-step Learning loss
    # we are gonna combine 1-step loss and n-step loss so as to
    # prevent high-variance. The original rainbow employs n-step loss only.
    if self.use_n_step:
        gamma = self.gamma ** self.n_step
        samples = self.memory_n.sample_batch_from_idxs(indices)
        elementwise_loss_n_loss = self._compute_dqn_loss(samples, gamma)
        elementwise_loss += elementwise_loss_n_loss

        # PER: importance sampling before average
        loss = torch.mean(elementwise_loss * weights)

    self.optimizer.zero_grad()
    loss.backward()
    clip_grad_norm_(self.dqn.parameters(), 10.0)
    self.optimizer.step()

    # PER: update priorities
    loss_for_prior = elementwise_loss.detach().cpu().numpy()
    new_priorities = loss_for_prior + self.prior_eps
    self.memory.update_priorities(indices, new_priorities)

    # NoisyNet: reset noise
    self.dqn.reset_noise()
    self.dqn_target.reset_noise()

    return loss.item()

```



```
def train(self, num_frames: int, plotting_interval: int = 200):
    """Train the agent."""
    self.is_test = False

    state = self.env.reset()
    update_cnt = 0
    losses = []
    scores = []
    score = 0

    for frame_idx in range(1, num_frames + 1):
        action = self.select_action(state)
        next_state, reward, done = self.step(action)

        state = next_state
        score += reward

        # NoisyNet: removed decrease of epsilon

        # PER: increase beta
        fraction = min(frame_idx / num_frames, 1.0)
        self.beta = self.beta + fraction * (1.0 - self.beta)

        # if episode ends
        if done:
            state = self.env.reset()
            scores.append(score)
            score = 0

        # if training is ready
        if len(self.memory) >= self.batch_size:
            loss = self.update_model()
            losses.append(loss)
            update_cnt += 1

            # if hard update is needed
            if update_cnt % self.target_update == 0:
                self._target_hard_update()

        # plotting
        if frame_idx % plotting_interval == 0:
            self._plot(frame_idx, scores, losses)

    self.env.close()

def test(self) -> List[np.ndarray]:
    """Test the agent."""
    self.is_test = True

    state = self.env.reset()
    done = False
    score = 0

    frames = []
    while not done:
```

```

frames.append(self.env.render(mode="rgb_array"))
action = self.select_action(state)
next_state, reward, done = self.step(action)

state = next_state
score += reward

print("score: ", score)
self.env.close()

return frames

def _compute_dqn_loss(self, samples: Dict[str, np.ndarray], gamma: float) -> torch.Tensor:
    """Return categorical dqn loss."""
    device = self.device # for shortening the following lines
    state = torch.FloatTensor(samples["obs"]).to(device)
    next_state = torch.FloatTensor(samples["next_obs"]).to(device)
    action = torch.LongTensor(samples["acts"]).to(device)
    reward = torch.FloatTensor(samples["rews"].reshape(-1, 1)).to(device)
    done = torch.FloatTensor(samples["done"].reshape(-1, 1)).to(device)

    # Categorical DQN algorithm
    delta_z = float(self.v_max - self.v_min) / (self.atom_size - 1)

    with torch.no_grad():
        # Double DQN
        next_action = self.dqn(next_state).argmax(1)
        next_dist = self.dqn_target.dist(next_state)
        next_dist = next_dist[range(self.batch_size), next_action]

        t_z = reward + (1 - done) * gamma * self.support
        t_z = t_z.clamp(min=self.v_min, max=self.v_max)
        b = (t_z - self.v_min) / delta_z
        l = b.floor().long()
        u = b.ceil().long()

        offset = (
            torch.linspace(
                0, (self.batch_size - 1) * self.atom_size, self.batch_size
            ).long()
            .unsqueeze(1)
            .expand(self.batch_size, self.atom_size)
            .to(self.device)
        )

        proj_dist = torch.zeros(next_dist.size(), device=self.device)
        proj_dist.view(-1).index_add_(
            0, (l + offset).view(-1), (next_dist * (u.float() - b)).view(-1)
        )
        proj_dist.view(-1).index_add_(
            0, (u + offset).view(-1), (next_dist * (b - l.float())).view(-1)
        )

        dist = self.dqn.dist(state)
        log_p = torch.log(dist[range(self.batch_size), action])

```

```

        elementwise_loss = -(proj_dist * log_p).sum(1)

    return elementwise_loss

def _target_hard_update(self):
    """Hard update: target <- local."""
    self.dqn_target.load_state_dict(self.dqn.state_dict())

def _plot(
    self,
    frame_idx: int,
    scores: List[float],
    losses: List[float],
):
    """Plot the training progresses."""
    clear_output(True)
    plt.figure(figsize=(20, 5))
    plt.subplot(131)
    plt.title('frame %s. score: %s' % (frame_idx, np.mean(scores[-10:])))
    plt.plot(scores)
    plt.subplot(132)
    plt.title('loss')
    plt.plot(losses)
    plt.show()

```

▼ Environment

You can see the [code](#) and [configurations](#) of CartPole-v0 from OpenAI's repository.

```

# environment
env_id = "CartPole-v0"
env = gym.make(env_id)
if IN_COLAB:
    env = gym.wrappers.Monitor(env, "videos", force=True)

```

▼ Set random seed

```

seed = 777

def seed_torch(seed):
    torch.manual_seed(seed)
    if torch.backends.cudnn.enabled:
        torch.backends.cudnn.benchmark = False
        torch.backends.cudnn.deterministic = True

np.random.seed(seed)
random.seed(seed)
seed_torch(seed)
env.seed(seed)

```

[777]

▼ Initialize



```
# parameters
num_frames = 20000
memory_size = 10000
batch_size = 128
target_update = 100

# train
agent = DQNAgent(env, memory_size, batch_size, target_update)

cpu
```

▼ Train

```
agent.train(num_frames)
```



▼ Test

Run the trained agent (1 episode).

```
frames = agent.test()

score: 200.0
```

▼ Render

```

if IN_COLAB: # for colab
    import base64
    import glob
    import io
    import os

    from IPython.display import HTML, display

def ipython_show_video(path: str) -> None:
    """Show a video at `path` within IPython Notebook."""
    if not os.path.isfile(path):
        raise NameError("Cannot access: {}".format(path))

    video = io.open(path, "r+b").read()
    encoded = base64.b64encode(video)

    display(HTML(
        data="""
        <video alt="test" controls>
        <source src="data:video/mp4;base64,{0}" type="video/mp4" />
        </video>
        """.format(encoded.decode("ascii"))
    ))

    list_of_files = glob.glob("videos/*.mp4")
    latest_file = max(list_of_files, key=os.path.getctime)
    print(latest_file)
    ipython_show_video(latest_file)

else: # for jupyter
    from matplotlib import animation
    from JSAnimation.IPython_display import display_animation
    from IPython.display import display

def display_frames_as_gif(frames):
    """Displays a list of frames as a gif, with controls."""
    patch = plt.imshow(frames[0])
    plt.axis('off')

    def animate(i):
        patch.set_data(frames[i])

    anim = animation.FuncAnimation(
        plt.gcf(), animate, frames = len(frames), interval=50
    )
    display(display_animation(anim, default_mode='loop'))

# display
display_frames_as_gif(frames)

```

videos/openaigym.video.0.886.video000125.mp4

0:00 / 0:04



▼ Conclusion

Thus in this Mini Project, we have observed the superior performance of Rainbow technique while solving the Cartpole problem and observed the performance of the model by rendering it on screen. It is interesting to observe the training curves and watching the model improve over time and plateau, with increase in training time. The end model is able to maintain the Cartpole balance as seen in the video above and we have learnt many interesting concepts about Reinforcement Learning in the process.

✓ 0s completed at 1:22 PM

